

X takes legal action against the Nitter project

X has sent cease-and-desist letters demanding the permanent shutdown of Nitter, an open source project that lets users read public X posts without logging into the platform or using its app. The letters, sent on August 24, 2026, reportedly also targeted multiple Nitter instances. Nitter.net is now offline and development has stopped temporarily while its creator, developer Zedeus, seeks legal advice.

A letter viewed by *TechCrunch* accuses Nitter of the ‘unlawful use and circumvention of X’s application programming interface (API) and associated data’, as well as scraping X data and accessing accounts and session tokens in violation of X’s rules.

Nitter said: “On 24 August 2026 cease and desist letters have been sent by X Corp. demanding a permanent takedown of Nitter instances and the project’s repository.”

The project fetched public X posts while removing ads, tracking cookies and JavaScript. Its code also powers services such as XCancel. In 2024, X API restrictions had already forced Nitter.net offline temporarily and required hosts to connect instances to real X accounts.

The latest action shifts the pressure from technical restrictions to legal enforcement against an open source project, its repository and independent instances. Nitter has said that it would respond to this action.



and remains an important part of Microsoft’s Windows App software developer’s kit (SDK) ecosystem.

Perplexity open sources AI inference engine Lily

Perplexity has open sourced Lily, a local artificial intelligence inference engine developed for Apple Silicon. The engine powers the local computing component of Hybrid Compute in Perplexity Computer and is designed specifically to run the Qwen3.6-35B-A3B model

Unlike general-purpose AI frameworks designed to support multiple models and hardware platforms, Lily takes a highly specialised approach. It uses a Rust runtime and handwritten Metal kernels optimised around the architecture of Qwen3.6-35B-A3B and Apple Silicon. Neither PyTorch nor Apple’s MLX framework is part of Lily’s execution path.

The engine provides an OpenAI-compatible chat-completions API for local text generation. Qwen3.6-35B-A3B is a mixture-of-experts model with 35 billion total parameters, although roughly 3 billion parameters are activated for each



token. Its architecture combines mixture-of-experts processing with grouped-query attention and Gated DeltaNet layers, creating different computational workloads that Lily specifically optimises for.

According to Perplexity’s reported benchmarks, Lily achieved an average 1.23 times higher prefill throughput and 1.35 times faster decoding than MLX-LM when tested across different context lengths on a 40-core, 128GB M5 Max system. The optimisations include keeping expert routing operations on the graphics processing unit (GPU) and combining weight dequantisation with grouped matrix operations to reduce memory overhead.

By open sourcing Lily, Perplexity is making its specialised local inference technology available to developers interested in running large AI models efficiently on Apple Silicon. The project demonstrates an alternative to general-purpose inference frameworks by focusing deeply on the requirements of a specific model and hardware architecture, while supporting Perplexity’s broader approach of combining local and cloud-based AI computing.

DeepSeek open sources its first native vision model

DeepSeek has open sourced V4-Flash-Vision-Exp, described as the company’s first native vision model. The experimental release expands DeepSeek’s artificial intelligence capabilities into multimodal computing, enabling the model to process and understand visual information alongside its broader AI functions.

A key part of the release is its open source availability. DeepSeek has made V4-Flash-Vision-Exp available through Hugging Face, giving developers and researchers a platform to access the model and experiment with its capabilities. The model is released under the MIT License, which permits broad use, modification and redistribution.

The experimental designation indicates that V4-Flash-Vision-Exp is part of DeepSeek’s continuing work on multimodal and next-generation AI technologies. The release provides an early look at the company’s approach to combining language and visual understanding within its AI ecosystem.

By releasing its first native vision model openly, DeepSeek is extending its open source AI strategy into multimodal systems. Availability through Hugging Face and the use of the MIT License allow the wider developer and research community to access, test, modify and potentially build new applications around the model’s vision capabilities.