

## Hugging Face launches Microduck robot

Hugging Face has unveiled Microduck, a 25cm-tall open source robot developed with its robotics subsidiary Pollen Robotics. The duck-shaped robot is designed as an accessible platform for experimenting with physical AI and reinforcement learning, allowing developers to teach it new behaviours and deploy them on the robot. Microduck can walk, crouch, recover from several common falls, pick up objects weighing up to 800 grams with its beak and even roller-skate.

The robot uses a camera, lidar sensors and two inertial measurement units (IMUs) to perceive its surroundings and monitor movement. Its compact design weighs less than 800 grams, making it small enough for desktop use while still allowing it to move through its environment. The design focuses more on physical interaction and movement than conversational interaction.

Microduck's software development kit, simulation environment and complete reinforcement-learning training stack are available on GitHub. Developers can train behaviours in simulation, transfer them to the physical robot, evaluate the results



and then retrain and redeploy the behaviour. This workflow is intended to make experimentation with physical AI more accessible and allow developers to share robotics behaviours in a manner similar to sharing AI models.

Pollen Robotics has positioned Microduck as a platform for learning how AI behaviours move from simulation into real-world machines. Rather than attempting to solve every robotics challenge, the project intends to provide developers, researchers and enthusiasts with a relatively accessible physical platform for experimenting with movement, reinforcement learning and robot control. The open software stack also gives users greater visibility into how behaviours are developed and deployed.

Microduck is available for pre-order at US\$ 399, before taxes and shipping, with first deliveries targeted before Christmas 2026 in North America, Europe and the UK.

## Arduino's VENTUNO Q board boosts edge AI

Arduino has opened pre-orders for the VENTUNO Q, a high-performance single-board computer designed to push the company deeper into edge artificial intelligence (AI), robotics and real-world automation. The board combines an eight-core Qualcomm processor with a dedicated microcontroller, creating a platform that handles both computationally intensive AI workloads and time-critical hardware control on the same system.

At the centre of the board is Qualcomm's Dragonwing IQ-8275 processor, featuring eight 64-bit Kryo central processing unit (CPU) cores running at up to 2.35GHz, an Adreno 623 graphics processing unit (GPU) and a Hexagon DSP-based AI accelerator.

## Nvidia adds vGPU support to open source Nova driver

Nvidia continues to expand the capabilities of its open source Nova kernel graphics driver for Linux. The company has posted a series of 13 patches introducing an Nvidia vGPU manager and a virtual function I/O (VFIO) variant driver, extending the functionality being developed around the upstream Nova driver.

The new patches are designed to enable vGPU management, allowing a physical GPU to support virtualised workloads. They also add support for VFIO users working with quick emulator (QEMU), enabling graphics processing unit (GPU) virtualisation for both Linux and Windows guest operating systems while using the upstream Nova kernel driver.

VFIO is commonly used in Linux virtualisation environments to provide virtual machines with more direct access to hardware. In this case, the new Nvidia work is aimed at integrating virtual GPU functionality into the Nova driver stack rather than relying solely on separate proprietary driver components.

However, the initial support is expected to focus primarily on Nvidia's workstation and data-centre GPUs, rather than consumer GeForce graphics cards. The patch series provides an early indication of how Nvidia plans to support virtualised GPU workloads as development of the open source Nova driver continues.

The Nova driver is still under development and is not yet ready for general end user deployment in the mainline Linux kernel. Nevertheless, the addition of vGPU management and VFIO support shows Nvidia continuing to build functionality around its open source Linux graphics driver, with future developments expected to expand its role in the company's Linux GPU software ecosystem.