




Figure 1: The role of the IPU in a server and data centre

operating systems, the raw muscle rendering immersive digital worlds, and the complex matrix-crunchers training the next generation of massive large language models (LLMs).

However, beneath the glamour and the staggering teraflops of computational power, there is a quiet, profound revolution occurring in the networking racks and server chassis of the modern data centre. Behind every highly publicized GPU cluster and every hyper-efficient TPU supercomputer, there is a vital layer of invisible infrastructure keeping the entire ecosystem from collapsing under its own weight.

Enter the infrastructure processing unit (IPU) — the unsung hero of the data centre. While CPUs, GPUs, and TPUs receive all the attention, IPUs are the dedicated workhorses managing the chaotic, high-volume flow of data, securing the perimeter, and ensuring that the high-profile processors can do exactly what they were designed to do: compute.

The hidden cost of the cloud: The ‘data centre tax’

To understand why IPUs are becoming the most critical component in hyperscale architecture, one must first understand the fundamental bottleneck of the modern cloud.

In a traditional, CPU-centric data centre, the central processing unit is the undisputed orchestrator of everything. Not only is it responsible for running the revenue-generating workloads, it is also burdened with managing the underlying infrastructure.

As cloud environments have evolved from monolithic applications to highly distributed, microservices-based

architectures, the volume of internal ‘East-West’ traffic within the data centre has exploded. To manage this, hyperscalers rely on software-defined networking (SDN), software-defined storage, and complex hypervisors. The problem is that running these virtualised infrastructure services requires a massive amount of computational power.

The infrastructure overhead tax: In a conventional cloud server, up to 30% to 50% of the CPU’s core capacity can be consumed by background infrastructure tasks. This includes managing virtual switches (vSwitches), routing network packets, encrypting and decrypting data, and managing storage input/output (I/O).

For a cloud service provider (CSP), this is a devastating inefficiency. If a server is equipped with a high-end 64-core processor, but 20 of those cores are busy just keeping the network and storage running, those are 20 cores that cannot be rented out to a paying customer. In an industry where revenue is directly tied to the utilisation of available compute, this ‘data centre tax’ represents billions of dollars in lost monetisation potential.

Decoding the IPU: What it is and how it works

An IPU is a highly advanced, programmable network device designed specifically to intelligently manage system-level infrastructure resources. It intercepts the data before it ever reaches the CPU, acting as a dedicated brain for the data centre’s plumbing.

While earlier iterations of network cards (like standard NICs) simply moved packets from a cable to the host memory,