

Feature	Traditional CPU-centric architecture	IPU-accelerated architecture
Workload placement	Infrastructure and tenant applications compete for the same CPU cores.	Clean separation: CPU handles tenant applications, IPU handles infrastructure.
Server monetisation	70% - 80% CPU utilisation available for billing.	100% CPU utilisation available for billing.
Security model	Software-based hypervisor isolation; vulnerable to host CPU exploits.	Hardware-based air-gap; infrastructure control plane is physically isolated.
Storage strategy	Localised, physical disks prone to stranded capacity.	Diskless servers utilising pooled, network-attached NVMe via the IPU.
Data throughput	Bottlenecked by CPU interrupts and software switching.	Line-rate hardware acceleration (400Gbps+) with zero CPU intervention.

Programmability: Moving at the speed of software

Hardware acceleration is not a new concept — application-specific integrated circuits (ASICs) have been doing it for decades. However, ASICs are famously inflexible. Once manufactured, their logic is permanently etched into the silicon. If a cloud provider wants to update a network protocol or deploy a new security encryption standard, an ASIC becomes obsolete.

The genius of the modern IPU lies in its programmability. By combining hardened, fixed-function accelerators (for tasks that will never change, like basic packet routing) with programmable cores and open software frameworks (such as P4 for data plane programming), IPU's offer the best of both worlds.

Cloud operators can customise their data pipelines, deploy advanced telemetry, and roll out new hyper-scale microservices through software updates that directly manipulate the IPU's behaviour. This allows data centres to adapt to shifting AI workloads, new cybersecurity threats, and evolving cloud orchestration frameworks without needing to physically replace the hardware in millions of servers.

IPUs: The backbone of the next-generation cloud

We have officially reached the physical limits of the

monolithic, CPU-does-it-all server architecture. As the world demands exponentially more compute power to fuel generative AI, real-time massive data analytics, and ultra-low-latency edge applications, the infrastructure must evolve to keep pace.

While the fascination will undoubtedly remain fixated on the raw intelligence of CPUs, the parallel processing power of GPUs, and the deep-learning mastery of TPUs, the true scalability of the cloud relies on an entirely different piece of silicon. IPU's are the dedicated guardians, the efficient traffic directors, and the tireless workers that allow the data centre to operate securely and profitably at a staggering, global scale.

The industry is waking up to this reality. The shift from compute-heavy nodes to heavily disaggregated, IPU-centric architectures is no longer just a theoretical roadmap — it is a financial and operational imperative. Because of their profound ability to reclaim revenue, harden security, and accelerate data pipelines, IPU's are incredibly important and are going to get a significant amount of investment from all cloud service providers in the next few years. **END** 🐧

References

- <https://community.intel.com/t5/Blogs/Tech-Innovation/Data-Center/The-IPU-A-New-Strategic-Resource-for-Cloud-Service-Providers/post/1335081#:~:text=Offloading%20infrastructure%20tasks%20to%20the,driven%20the%20development%20of%20IPUs>
- <https://simplyblock.io/glossary/infrastructure-processing-unit/#:~:text=for%20higher%20efficiency.%,How%20an%20IPU%20Handles%20Infrastructure%20Workloads,traffic%20routing%2C%20and%20tenant%20isolation>
- <https://www.presidio.com/blogs/intel-ipus-and-smartnics-boost-data-center-utilization-and-lower-tco/#:~:text=Introduced%20in%20June%202021%2C%20the,intelligently%20accelerates%20and%20manages%20them>
- <https://medium.com/intel-tech/the-infrastructure-processing-unit-for-the-next-generation-cloud-services-3d813d1d4ed4>
- <https://www.youtube.com/watch?is=pXjE-x2L8ipNjSN8&v=6Z20LabHy48&feature=youtu.be>
- <https://www.trentonsystems.com/en-us/resource-hub/blog/what-is-an-ipu#:~:text=It%20allows%20cloud%20operators%20to,architecture%20in%20the%20data%20center>
- <https://semiwiki.com/semiconductor-manufacturers/intel/361110-intels-ipu-e2200-redefining-data-center-infrastructure/#:~:text=It%20also%20plays%20a%20crucial,Summary>

By: Ishaan Gupta

The author leads a team working on compute accelerators for AI workloads at Google. In his 13+ years of experience in the software and AI industry, he has worked in the augmented reality AI glasses organisation at Meta, and in Square, the payments company.