

The Benefits and Uses of Open Source LLMs

There are several open source alternatives available today for expensive commercial large language models. Engineering teams can easily leverage open-weight models to achieve tangible ROI across production environments.



Akki Bharath Goud, cloud architect, OVHcloud

move into production without a clear ROI, organisations need accessible AI tools. The easiest way to consume AI without massive training investments is through open source LLMs, as thousands are readily available for research and deployment in public repositories/platforms like Hugging Face.

Hardware requirements and infrastructure strategy

Regarding infrastructure and usage, there are multiple ways to leverage off-the-shelf, pre-trained models. One approach is to download the model weights directly. You must evaluate the required hardware resources -- some large language models have high memory demands or specific GPU requirements, such as an H100. Models such as gpt-oss-120b are effective for reasoning and agentic tasks. Platforms such as Hugging Face allow you to research model architectures, review code snippets, and carry out cost comparisons.

Infrastructure providers such as OVHcloud host open source models and offer ready-to-use inference endpoints built on their public cloud infrastructure, featuring top-performing models selected directly from open source repositories. OVHcloud's public cloud is built on the open source OpenStack platform, reflecting its commitment to technology sovereignty. The key advantage is that you pay strictly for inference costs rather than for model licensing. A full list of available endpoints can be found in OVHcloud's AI Endpoints catalogue.

Let's begin with one important question: why do we need open source LLMs?

The answer: we have well-trained open source LLMs that are trained on huge datasets, whose weights the community has publicly shared. They possess tremendous knowledge (whether they have 70 billion parameters, 20 billion parameters, or another scale) acquired through training processes that utilise significant GPU resources. Because these are available on public repositories/platforms, you can download, utilise, and reuse them without paying extra fees beyond your inference costs. While commercial models exist, open source LLMs can perform these same jobs effectively.

The evolution of enterprise AI and market trends

Earlier, AI was often treated as an experimental phase where every department was asked to start at least one AI project, regardless of its ultimate utility. That trend has changed; leaders now focus on whether an AI project aligns with their broader strategic goals.

Furthermore, the technology is moving very quickly. In 2025, the focus was on genAI applications, LLMs, and building RAG systems. Moving into 2026, the discussion has shifted towards AI agents and agentic workflows, and autonomous agents capable of managing lower-end operational tasks. Because proof-of-concept projects no longer