

```

- name: models-from-internal-registry-only
  match:
    any:
      - resources:
          kinds: ["Pod"]
          namespaces: ["inference-*"]
  validate:
    message: "Model artifacts must come from the internal registry."
  pattern:
    spec:
      containers:
        - image: "registry.internal.example.com/*"
        env:

```

```

- name: MODEL_URI
  value: "s3://models-approved/*"

```

Useful companion policies are: deny pods in inference namespaces without a cost-attribution label, require a resource claim rather than a bare GPU count, block egress-enabled tool sandboxes outside an allowlisted namespace, and require a signature attestation on any image that loads model weights.

How do you control the cost of inference?

Inference cost scales with usage and lands in gross margin. Most teams reach for model selection first because it feels like the engineering lever, then find that caching and prompt structure return more, sooner, with less risk.




Figure 5: A routing cascade. Escalation on failure keeps quality intact while making the expensive model the exception rather than the default