

what an error message really means, draft infrastructure configurations, or come up with solutions that look plausible. Of course, the engineer still has the final say — reviewing the output and executing it.

In the next stage, AI will be able to carry out several connected tasks, with only slight oversight. For instance, an agent may be given the goal of investigating a deployment that failed. Instead of just explaining the error, it could gather the necessary details, reason through the failure, surface a few likely causes, and then put together a remediation plan.

Human-in-the-Loop

The idea of autonomous systems often brings concerns about safety and control, and honestly, that matters more than people think. In DevOps, these worries become especially important, because a wrong move can hit production systems, mess with customer data, weaken security, and even disrupt business operations.

An agent may be allowed to restart a development service on its own, but a change in the production infrastructure may still need a human sign-off. Likewise, an agent can automatically investigate a security alert, but it would pause for authorisation before it touches a firewall rule, or a production access policy, depending on the scope.

So you end up with this Human-in-the-Loop model, where people and agents collaborate. The AI deals with repetitive investigation and operational work, while engineers make the high impact calls.

The goal isn't to remove humans from DevOps at all. It's more to shift human attention towards the areas where judgment counts the most.

The rise of the multi-agent DevOps team

Agentic DevOps gets even more powerful when several specialised agents team up. A release agent

obsesses over deployments and CI/CD, making sure version changes don't slip. An infrastructure agent handles the cloud side and Kubernetes objects, all the way from namespaces to scaling decisions. The observability agent watches logs, metrics, and traces, looking for weird patterns. A security agent goes after vulnerabilities and policy violations that look small but aren't. And an SRE agent zeroes in on reliability plus incident response when things start acting up.

Picture a production incident where application latency suddenly spikes. The observability agent catches the anomaly first, and then the deployment agent checks whether a recent release was involved. The infrastructure agent digs into resource utilisation, comparing current load with the baseline. Meanwhile, the security agent looks for suspicious activity, connections, or any strange authorisation behaviour. After that a coordinating agent bundles everything together, synthesizes the signals, and recommends what to do next.

This setup pushes DevOps closer to the idea of an AI-powered engineering team, not just automation for automation's sake.

What agentic DevOps means for engineers

The emergence of agentic DevOps will also nudge the role of DevOps engineers. A lot of the repetitive operations work that used to pile up can increasingly be assisted or even handled by intelligent systems. Engineers may find themselves spending less time manually digging

through logs, running routine commands, crafting repeat configs, and following common alert storms.

At the same time though, fresh duties will show up. Engineers will need to grasp how to design agentic workflows, equip the agents with the right tools, set up permissions properly, assess what the agents decide, keep an eye on how they behave, and build solid safety mechanisms that hold up when things get weird.

The road to autonomous DevOps

Agentic DevOps should not be treated like some sudden, full stop replacement for what we already do in DevOps. It is more like an evolution stacked on top of existing practices, and not a clean swap. CI/CD, Infrastructure as Code, cloud platforms, Kubernetes, observability, GitOps and security all stay non-negotiable. What AI agents add is a new intelligent layer that can interact with these pieces in a more proactive way.

You can think of the path as going from manual operations to automation, then to automation with smarter help, and finally towards something closer to controlled autonomy.

The end game is a software delivery and operations environment where systems can notice problems, grasp the surrounding context, pick a fitting response, perform only the actions that are approved, and then confirm the outcome.

Agentic DevOps is not only about fitting AI onto a pipeline. It is a wider shift in how software platforms get designed, and how they run day to day. **END** 🐼

By: D. Banerjee and Sourav K.

D. Banerjee is an IT professional with 15+ years of experience in Tier-1 IT companies, specialising in delivering scalable enterprise solutions and leading complex technology initiatives.

Sourav K. is an IT professional with around 5 years of experience in Tier-1 IT companies, skilled in modern technologies and focused on building efficient, data-driven solutions.